id	what happened	the lab analogue	test?
T01	The gate read pytest exit 1 as 'the test failed' when the test never ran (pytest missing, nothing collected, a syntax error).	A QC step marked FAIL when the instrument never ran — the sample is discarded for a reason that has nothing to do with the	yes
T02	Trace publishing printed a URL and exited 0 while the query store recorded nothing.	An instrument log that says 'saved' while the LIMS holds no record; the run is unrecoverable at analysis time.	yes
T03	The lessons pin was computed exactly and identified nothing: a tree SHA over a store that background hooks rewrite.	A reagent lot number copied from the wrong bottle — exact, logged, and not the lot that was used.	yes
T04	The close commit swept hook-mutated memory files into build-unit commits (git add -A), so a unit's diff no longer described the unit.	A sample contaminated during handling: the assay result is real, but it is not the sample's alone.	yes
T05	`instinct capture` on an existing name discarded the correction and reinforced the stale text, exit 0.	A corrected lab-notebook entry overwritten by the original, with the correction counted as confirmation.	yes
T06	Run records could not say whether a run happened; a seeded fixture in the same schema was indistinguishable from a real record.	A results table with no provenance column — simulated and measured rows look the same.	yes
T07	`register verify-record` returned 1 for UNVERIFIED, and 1 means 'retry the unit' everywhere else in the tool.	Two assays sharing one flag value with opposite meanings; the downstream rule fires on the wrong one.	yes
T08	The hard task closed first try because the implementer prompt named the answer; nothing pins the prompts.	A positive control that answers the experimental question for you.	no
T09	Two build units were not testable in isolation; a park for that reason would have been indistinguishable from one earned by a hard task.	An assay whose negative result cannot separate reagent failure from a true negative.	no
T10	Test isolation depended on pytest's collection order: stubs installed at import time leaked for the whole session (F08).	Benchmark leakage — a held-out set the model has already seen, giving a score that measures the leak.	yes
T11	Five of the guard's own tests could not fail; two of three enforcement halves were deletable with the suite green; the mutation gate credited any non-zero exit as a kill (pass 5).	A held-out metric that stays good after the model is ablated — the evaluation is not measuring the model.	yes
T12	A pre-imported name was a transitive dependency: when the resolution changed, three tests went red (A7b).	An undeclared reagent that arrived in another kit's box; the protocol fails the day the kit changes.	yes
T13	The demo seeder wrote a made-up 23-task, 3-pass fixture into the live run-state directory, over pass 1's own record.	A fixture presented as data — a figure drawn from simulated rows placed in the results directory.	yes
T14	The referee's normaliser recognised two verdicts, so a skipped test read as a pass.	A skipped sample counted as a pass because the plate map had a well for it.	yes
T15	The register and the referee parsed the same test report differently, so one could pass what the other refused.	Two readers of one gel using different thresholds; the band is 'present' or not depending on who looked.	yes
T16	A suite with no declared count: a result the parser did not recognise was dropped, and a dropped result is indistinguishable from a test that never existed.	Missing samples not counted as missing — the denominator shrinks to what was measured.	yes
T17	`referee run --sign` attested that a run happened, not that it passed: a failing suite signed the same receipt as a green one.	A certificate of analysis that certifies only that an analysis was run.	yes
T18	Receipts issued before the fix still attested the property nobody had checked; receipt_version had to be bumped to 2.	Recalling a batch of certificates after the assay behind them is corrected.	yes
T19	`plan_required: true` was rendered 'True' by the config layer, the string compare failed, and the guard turned itself off — for the one consumer who asked for it.	A safety interlock that disengages when set to ON because the switch reports 'ON' in a different case.	yes
T20	A green gate could ship with zero tests executed: an empty test command was 'skip', and the DoD attested around it.	A QC report with the QC step left blank and the batch released anyway.	yes
T21	The referee command's 'could not run' (unknown commit, missing tool, zero tests collected) would have looked like 'the suite failed' — the draft wrapped pytest in `|| true`.	A negative result recorded when the instrument was off — the same channel carries both.	yes
T22	The evidence recorder read the referee's LAST output line and, under --sign, found no JSON: a signed, passing 75/75 run was recorded as NOT MEASURED.	A positive result filed as inconclusive because the parser looked at the footer, not the table.	yes